

# Boosting - AdaBoost

Zhe Yang



# Contents

## ❖ Boosting Introduction

## ❖ AdaBoost

- Algorithm
- Additive structure
- Different view of AdaBoost
- Investigate the loss function
- Margin theory

# Boosting

❖ Motivation: “Can a set of weak learners create a single strong learner?”

- Kearns & Valiant (1988, 1989)
- Weak learner: a ML algorithm with performance slightly better than random guessing

Yes !




❖ Brief history

- 1990 – First simple boosting procedure (Schapire)
  - Showed that a weak learner could always improve its performance
  - Initial classifier + two classifiers trained on misidentified/disagreed data by previous classifier
- 1993 – Idea of “Boost by majority” (Freund)
- ---- Border between theoretical and practical ----
- 1995 – AdaBoost (Freund & Schapire)
  - Significant improvement, combines many weighted weak learners, adjusted training weight

# AdaBoost

## Use binary classification as an example



**FIGURE 10.1.** Schematic of AdaBoost. Classifiers are trained on weighted versions of the dataset, and then combined to produce a final prediction.

---

### Algorithm 10.1 *AdaBoost.M1.*

---

1. Initialize the observation weights  $w_i = 1/N$ ,  $i = 1, 2, \dots, N$ .
2. For  $m = 1$  to  $M$ :
  - (a) Fit a classifier  $G_m(x)$  to the training data using weights  $w_i$ .
  - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
  - (c) Compute  $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$ .
  - (d) Set  $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$ ,  $i = 1, 2, \dots, N$ .
3. Output  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .

---

#### Points

- It's a **sequential** process
  - Updated weights on training sample, focusing on mistakes made by previous classifiers
- It's a family of **ensemble** technique
  - Weighted combination of classifiers at different stage, good classifier's decision is more important

#### Details

- Base classifier  $G$  is not specified
  - Decision tree is an ideal choice

# Why AdaBoost has big impact

## ❖ It's simple yet effective


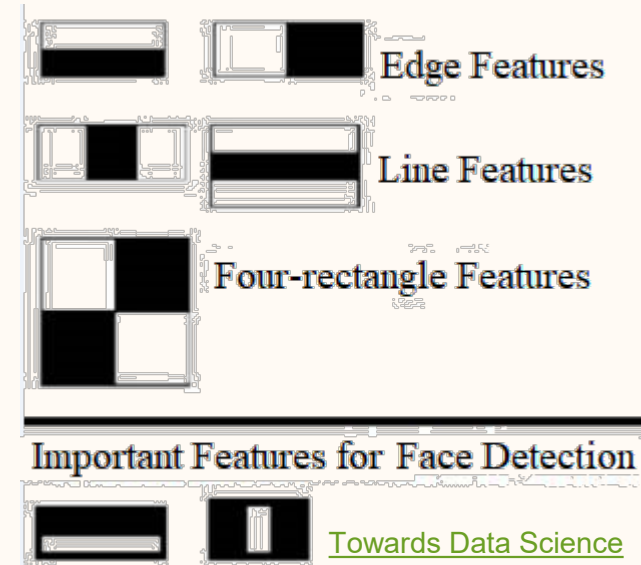
- The idea behind can be applied to different fields
- Computer vision: Viola-Jones detector
  - Base classifiers are the [haar-like feature](#) with examples given on the right side
  - On average only 8 features need to be evaluated



$$F(x) = \alpha_1 f_1(x) + \alpha_2 f_2(x) + \alpha_3 f_3(x) + \dots$$

## ❖ Guaranteed by the theory to be valid

- Train err decrease exponentially
- “Never” overtrained in practice



# AdaBoost: additive structure

- ❖ The structure of AdaBoost is indicated in the expression of final prediction

$$G(x) = \text{sign}\left(\sum_{m=1}^M \alpha_m G_m(x)\right)$$

- The basis functions are the individual classifiers  $G_m(x) \in \{-1, 1\}$

- ❖ More generally, basis function expansions takes the form

$$f(x) = \sum_{m=1}^M \beta_m b(x; \gamma_m)$$

- $\beta_m, m = 1, 2, \dots, M$  are the **expansion coefficients**
- $b(x; \gamma_m) \in \mathbb{R}$  are usually **simple functions**
  - Taking multivariate **input**  $x$
  - Defined by **parameter set**  $\gamma$

# Use of additive expansions

- ❖ Additive expansions are at the heart of many of the learning techniques
  - Neural network with single hidden layer
    - $b(x; \gamma) = \sigma(\gamma^T x)$ ,  $\sigma$  is the activation function,  $\gamma$  are the linear combination weights for input  $x$
  - Multivariate adaptive regression splines (MARS)
    - $b$  is from truncated-power spline basis functions,  $\gamma$  is the choice of  $x_{ij}, i = 1, 2, \dots, N, j = 1, 2, \dots, p$
  - Tree methods
    - $b$  is simple tree,  $\gamma$  defines splits and prediction at leaf nodes

# Forward stagewise additive modeling

- ❖ The additive models discussed in previous page are typically fit by minimizing a loss function averaged over the training data

$$\min_{\{\beta_m, \gamma_m\}_i^M} \sum_{i=1}^N L(y_i, \sum_{m=1}^M \beta_m b(x_i; \gamma_m))$$

- The loss function  $L$  could be the squared-error or a likelihood-based loss function
- ❖ It requires computationally intensive numerical optimization techniques to optimize with **all  $M$**  basis functions considered **simultaneously**, a simple alternative often can be found when it is feasible to rapidly solve the subproblem of fitting just a **single** basis function

$$\min_{\{\beta, \gamma\}} \sum_{i=1}^N L(y_i, \beta b(x_i; \gamma))$$



# Forward stagewise additive modeling (Algorithm)

---

**Algorithm 10.2** *Forward Stagewise Additive Modeling.*

---

1. Initialize  $f_0(x) = 0$ .

2. For  $m = 1$  to  $M$ :

(a) Compute

$$(\beta_m, \gamma_m) = \arg \min_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \beta b(x_i; \gamma)).$$

(b) Set  $f_m(x) = f_{m-1}(x) + \beta_m b(x; \gamma_m)$ .

---

At each stage  $m$  only update  $\beta_m$  and  $\gamma_m$

- by fitting the **partial residual**  $y - f_{m-1}(x)$
- **without adjusting** those have been already fixed in previous stage  $1, 2, \dots, m-1$
- we show next that AdaBoost is doing the same **equivalently**
  - even though we didn't use any loss function in the AdaBoost algorithm described earlier

# AdaBoost as a FSAM

- ❖ AdaBoost is equivalent to forward stagewise additive modeling (FSAM) using the loss function

$$L(y, f(x)) = \exp(-yf(x))$$

- The reason to use exponential form discussed later
- Note that we have  $f(x) \in \{-1, 1\} \rightarrow yf(x) = \begin{cases} 1, y = f(x) \\ -1, y \neq f(x) \end{cases}$
- ❖ At stage  $m$ , the optimization is done by solving

$$(\beta_m, G_m) = \arg \min_{\beta, G} \sum_{i=1}^N \exp[-y_i(f_{m-1}(x_i) + \beta G(x_i))]$$

- ❖ Rewrite with  $w_i^m = \exp(-y_i f_{m-1}(x_i))$

$$(\beta_m, G_m) = \arg \min_{\beta, G} \sum_{i=1}^N w_i^m \exp[-y_i \beta G(x_i)]$$

- $w_i^m$  is **fixed** in previous stages and can be regarded as a weight applied to observations
- Note  $G_m(x)$  also have  $G_m(x) \in \{-1, 1\} \rightarrow yG_m(x) = \begin{cases} 1, y = G_m(x) \\ -1, y \neq G_m(x) \end{cases}$

# Get $(\beta_m, G_m)$

- ❖  $\beta_m$  is a positive constant even though it's unknown yet, we consider it as fixed and solve the  $G_m$

$$G_m = \arg \min_G \sum_{i=1}^N w_i^{(m)} I(y_i \neq G(x_i))$$

- we only consider  $y_i \neq G(x_i)$  situation here, because the  $y_i = G(x_i)$  situation is the other side of the coin

$$I(y_i = G(x_i)) = 1 - I(y_i \neq G(x_i))$$

- ❖ With known  $G_m$  the loss function  $L_m(\beta)$  will be

$$L_m(\beta) = \sum_{i=1}^N w_i^{(m)} \exp(-\beta y_i G_m(x_i)) = e^{-\beta} \sum_{y_i = G_m(x_i)} w_i^{(m)} + e^{\beta} \sum_{y_i \neq G_m(x_i)} w_i^{(m)}$$

- ❖ Get optimal  $\beta$  to minimize  $L_m(\beta)$

$$\beta_m = \frac{1}{2} \ln \left( \frac{1 - err_m}{err_m} \right), err_m = \frac{\sum_{i=1}^N w_i^{(m)} I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i^{(m)}}$$

- Note that  $\beta_m$  could be **negative** if  $err_m < 0.5$ , in which case it automatically reverse the polarity.
  - Move the “-1” in  $\beta_m$  to  $G_m$ :  $\beta_m \rightarrow -\beta_m > 0$ ,  $G_m(x) \rightarrow -G_m(x)$ ,  $err_m \rightarrow 1 - err_m > 0.5$

# Stage updates

❖ In each stage  $m$ , the prediction function is updated as

$$f_m(x) = f_{m-1}(x) + \beta_m G_m(x)$$

- where  $\beta_m$  and  $G_m$  are estimated as illustrated in the previous page

❖ In next stage  $m + 1$

- the weight  $w_i^{(m)}$  in loss function is updated to  $w_i^{(m+1)}$  by

$$w_i^{(m+1)} = w_i^{(m)} \cdot e^{-\beta_m y_i G_m(x_i)}$$

- then  $\beta_{m+1}$  and  $G_{m+1}$  are updated following the same procedures

# AdaBoost as a Newton-like method

- ❖ Again, start from the exponential loss  $L(y, f(x)) = \exp(-yf(x))$ . Suppose we seek an improvement

$$f(x) + cg(x)$$

- here  $c$  is a constant to be determined and  $g(x) \in \{-1, 1\}$  giving a binary prediction
- ❖ We can expand  $L$  at  $g(x) = 0$

$$\begin{aligned} L(y, f(x) + cg(x)) &= \exp(-yf(x)) \exp(-ycg(x)) \\ &\approx \exp(-yf(x)) \left[ 1 - cyg(x) + \frac{c^2}{2} \right] \stackrel{\text{def}}{=} L' \end{aligned}$$

- notice the derivation of  $\frac{c^2}{2}$  use the fact that  $y^2 = 1$  and  $g(x)^2 = 1$
- $\exp(-yf(x))$  is determined by previous process and can be view as weight  $w$

# Get $g(x)$

- ❖ Minimizing  $L'$  with  $c$  fixed is equivalent to maximizing  $yg(x)$ , the solution is

$$g(x) = \begin{cases} 1, & p(y = 1) > p(y = -1) \\ -1, & \text{otherwise} \end{cases}$$

- ❖ Again, use the fact that  $y^2 = 1$  and  $g(x)^2 = 1$

$$yg(x) = -\frac{[y - g(x)]^2}{2} + 1$$

- ❖ Then minimizing  $L'$  is equivalent to minimizing the least-square loss

$$\frac{[y - g(x)]^2}{2} - 1$$

# Get $c$

- ❖ With solved  $g(x) \in \{-1, 1\}$ , we can get  $c$  again by optimizing  $L$ . The essential component is

$$\exp(-ycg(x)) = \sum_{y_i=g(x_i)} e^{-c} + \sum_{y_i \neq g(x_i)} e^c = n_T e^{-c} + n_F e^c \stackrel{\text{def}}{=} l(c)$$

- where  $n_T$  is the number of correct classification and  $n_F$  is the number of incorrect classification
- ❖ Minimizing the equation above

$$\frac{\partial l(c)}{\partial c} = -n_T e^{-c} + n_F e^c = 0 \rightarrow c = \frac{1}{2} \ln \frac{n_T}{n_F} = \frac{1}{2} \ln \frac{1 - \text{err}}{\text{err}}$$

- where  $\text{err} = \frac{n_F}{n_T + n_F} = \frac{n_F}{n}$

# Stage updates

❖ We get the exactly same form of coefficient and update function as what we got when we consider AdaBoost as FSAM

❖ And we can now construct a Newton-like method's update in each stage

$$f(x) \rightarrow f(x) + \frac{1}{2} \ln \frac{1 - err}{err} g(x)$$

❖ In next round, the weight  $w$  is updated by

$$w \rightarrow w \cdot e^{-cyg(x)} = w \cdot e^{-c(1-2I(y \neq g(x)))}$$



# What does exponential loss estimate?

- ❖ Investigating the population minimizer of  $\exp(-yf(x))$

$$f^*(x) = \arg \min_{f(x)} E_{y|x}(e^{-yf(x)})$$

- Get derivative  $E_{y|x}(e^{-yf(x)})$  w.r.t  $f(x)$  and set to zero

$$\begin{aligned} \frac{\partial E_{y|x}(e^{-yf(x)})}{\partial f(x)} &= E_{y|x}(-ye^{-yf(x)}) = E_{y=1|x}(-ye^{-yf(x)}) + E_{y=-1|x}(-ye^{-yf(x)}) \\ &= -\Pr(y = 1|x) e^{-f(x)} + \Pr(y = -1|x) e^{f(x)} = 0 \end{aligned}$$

- Solve the equation

$$f^*(x) = \frac{1}{2} \ln \left( \frac{\Pr(y = 1|x)}{\Pr(y = -1|x)} \right) \text{ or } \Pr(y = 1|x) = \frac{1}{1 + e^{-2f^*(x)}}$$

- ❖ It is estimating the log-odds of  $\Pr(y = 1|x)$

- This provides an inspect to justify that we use the sign of the model prediction as classification  $\{-1, 1\}$ , where we have  $G(x) = \text{sign}(\sum_{m=1}^M \alpha_m G_m(x))$ .
  - if we have higher probability to predict  $y = 1$ ,  $f^*(x)$  will be “+”
  - if we have higher probability to predict  $y = -1$ ,  $f^*(x)$  will be “-”,

# Binomial deviance could be another candidate

- ❖ We now show that binomial deviance (cross-entropy) leads to same solution at the **population level**. Using  $y' = \frac{y+1}{2}$  to map  $\{-1, 1\}$  to  $\{0, 1\}$ , we have the binomial deviance loss

$$l(y, p(x)) = y' \log p(x) + (1 - y') \log(1 - p(x))$$

- ❖ Get derivative  $E_{y|x}(l(y, p(x)))$  w.r.t  $p(x)$  and set to zero

$$\begin{aligned} \frac{\partial E_{y|x} l(y, p(x))}{\partial p(x)} &= E_{y=1|x} \left( \frac{1}{p(x)} \right) + E_{y=-1|x} \left( -\frac{1}{1-p(x)} \right) \\ &= \frac{\Pr(y = 1|x)}{p(x)} - \frac{\Pr(y = -1|x)}{1-p(x)} = 0 \end{aligned}$$

- ❖ Solve the equation we get

$$p(x) = \Pr(y = 1|x) = \frac{1}{1 + e^{-2f(x)}}$$

- ❖ Once  $p(x)$  is determined,  $f(x)$  is also determined via a logit relation

# Continue

- ❖ Use  $f(x)$  to replace  $p(x)$

$$l(y, p(x)) = y' \log p(x) + (1 - y') \log(1 - p(x)) = \log \frac{e^{(y-1)f(x)}}{1 + e^{-2f(x)}}$$

- Notice that  $y \in \{-1, 1\}$

$$l(y, p(x)) = \begin{cases} \log \frac{1}{1 + e^{-2f(x)}}, y = 1 \\ \log \frac{e^{-2f(x)}}{1 + e^{-2f(x)}}, y = -1 \end{cases} = -\log(1 + e^{-2yf(x)})$$

- Expand  $e^{-yf(x)}$  and  $-l(y, p(x))$  to Taylor series, we have

$$-l(y, p(x)) = \log(2) - yf(x) + \frac{y^2}{2} f^2(x) + \dots$$

$$e^{-yf(x)} = 1 - yf(x) + \frac{y^2}{2} f^2(x) + \dots$$

- ❖ Binomial deviance loss and exponential loss are equivalent to second order

$$-l(y, p(x)) \approx e^{-yf(x)} + \log(2) - 1$$

# Compare different loss for classification


❖ In the right plot, “ $yf$ ” serves as a **margin** which plays a role analogous to the residuals  $y - f(x)$

❖ Investigate positive  $yf$

- Squared Error will penalize correct predictions
- All other loss decrease monotonically

❖ Investigate negative  $yf$

- All the loss penalize incorrect predictions but with different degree
- Exponential loss will penalize much more on incorrect predictions so it's less robust
  - If inputs have mislabeled entries, the performance of AdaBoost will degrade dramatically (empirical observation)




# AdaBoost doesn't overfit

- ❖ AdaBoost often doesn't overfit in practice
- ❖ Freund & Schapire[JCSS97] proved that the generalization error of AdaBoost is bounded by:

$$\epsilon_D \leq \epsilon_D + \tilde{O} \left( \sqrt{\frac{dM}{n}} \right)$$

- with probability at least  $1 - \delta$ , where  $d$  represents the complexity of base learners,  $n$  is the number of training instances,  $M$  is the number of learning rounds and  $\tilde{O}$  hides the logarithmic and constant terms
- It implies that AdaBoost will **overfit** if  $M$  is large
- ❖ Is there **tighter** bound to explain the observation in practice?

A typical performance plot of AdaBoost on real data



Seems contradict with the **Occam's Razor**

# Margin

❖ We have discussed the statistical view above. However, it did not explain why **AdaBoost is resistant to overfitting** according to observation in practice.

■ This can be explained by using **margin**

## ❖ Example


■ Binary classification can be viewed as the task of separating classes in a feature space

■  $f(x) = \text{sign}(w^T x + b)$

■ Margin of a single classifier  $f$ :  $yf(x)$

■ Margin of the additive model  $F$

$$yF(x) = \sum_{m=1}^M \alpha_m y f_m(x)$$



The bigger the margin,  
the higher the predictive confidence

# Margin explanation of AdaBoost


- ❖ Based on the concept of margin, Schapire et al. [1998] proved that, given any threshold  $\theta > 0$  of margin over the training data  $D$ , with probability at least  $1 - \delta$ , the test error of the ensemble  $\epsilon_D = P_{x \sim D}(y \neq F(x))$  is bounded by

$$\begin{aligned}\epsilon_D &\leq P_{x \sim D}(yF(x) \leq \theta) + \tilde{O}\left(\sqrt{\frac{d}{n\theta^2} + \ln \frac{1}{\delta}}\right) \\ &\leq 2^M \prod_{m=1}^M \sqrt{\epsilon_m^{1-\theta} (1 - \epsilon_m)^{1+\theta}} + \tilde{O}\left(\sqrt{\frac{d}{n\theta^2} + \ln \frac{1}{\delta}}\right)\end{aligned}$$

- ❖ This bound implies that, when other variables are fixed, the larger the margin over the training data, the smaller the test error

# Margin explanation of AdaBoost (continued)

- ❖ Why AdaBoost tends to be resistant to overfitting?
- ❖ Margin theory: Because it is able to increase the ensemble margin even after the training error reaches zero



- This explanation is quite intuitive
- It receives good support in empirical study



# The minimal margin bound

❖ Schapire et al.'s bound depends heavily on the smallest margin

■ If the smallest margin is large,  $P_{x \sim D}(yF(x) \leq \theta)$  will be small

❖ Thus, by considering the minimum margin:

$$\varrho = \min_{x \in D} yF(x)$$

❖ Breiman [Neural Comp. 1999] proved a test error bound ( $O\left(\frac{\log n}{n}\right)$ ), which is tighter than Schapire et al.'s bound


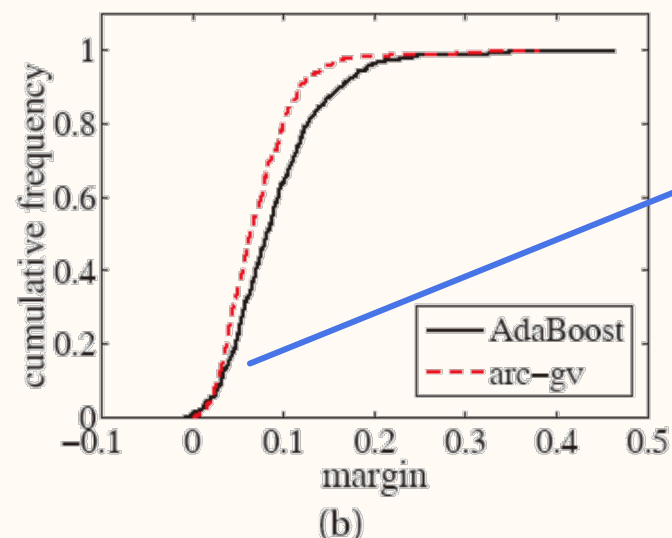

$$\epsilon_D \leq R \left( \ln(2n) + \ln\left(\frac{1}{R}\right) + 1 \right) + \frac{1}{n} \ln \frac{|H|}{\delta}$$

# The doubt about margin theory

- ❖ Breiman [Neural Comp. 1999] designed a variant of AdaBoost, the arc-gv algorithm, which directly maximizes the minimum margin
  - the margin theory would appear to predict that arc-gv should perform better than AdaBoost
- ❖ However, experiments show that, comparing with AdaBoost:
  - arc-gv does produce **uniformly larger minimum margin**
  - the **test error increases drastically** in almost every case
- ❖ Thus, Breiman convincingly concluded that **the margin theory was in serious doubt**. This almost sentenced the margin theory to death

# 7 years later

- ❖ Reyzin & Schapire [ICML'06 best paper] found that, amazingly, Breiman had not controlled model complexity well in exps
  - Breiman controlled the model complexity by using decision trees with a **fixed number of leaves**
  - Reyzin & Schapire found that, the trees of arc-gv are generally “**deeper**” than the trees of AdaBoost
- ❖ Reyzin & Schapire repeated Breiman's exps using decision stumps with two leaves: arc-gv is with larger minimum margin, but worse margin distribution
  - R&S claimed that the minimum margin is not crucial, and the **average** or **median** margin is crucial



# Supporting theory

## ❖ Equilibrium margin (Emargin) bound

- Considered factors different from Schapire et al. and Breiman's bounds
- No intuition to optimize

## ❖ The kth margin bound

- The minimum margin bound and Emargin bound are special cases of the kth margin bound, both are single-margin bound (not margin distribution bound)

## ❖ Finally

- Random choice of sample  $S$  with size  $n \geq 5$ , every voting classifier  $f \in \mathcal{C}(H)$  satisfies

$$\epsilon_{\mathcal{D}} \leq \frac{2}{n} + \inf_{\theta \in (0,1]} \left[ \Pr_S[yf(x) < \theta] + \frac{7\mu + 3\sqrt{3\mu}}{3\mu} + \sqrt{\frac{3\mu}{n} \Pr_S[yf(x) < \theta]} \right]$$

$O(\log n / n)$

- Where

$$\mu = \frac{8}{\theta^2} \ln n \ln(2|H|) + \ln \frac{2|H|}{\delta}.$$

- Uniformly tighter than Breiman's as well as Schapire et al.'s bounds
- Considers the same factors as Schapire et al. and Breiman
- thus, defends the margin theory against Breiman's doubt

# Variance of margin also matters

- ❖ We should pay attention to both average margin and margin variance.

$$\epsilon_{\mathcal{D}} \leq \frac{1}{n^{50}} + \inf_{\theta \in (0,1]} [\Pr_S[yf(x) < 0] + n^{-\frac{2}{1 - E_S^2[yf(x)] + \frac{\theta}{9}}} + \frac{3\sqrt{\mu}}{n^{\frac{3}{2}}} + \frac{7\mu}{3n} + \sqrt{\frac{3\mu}{n} \hat{\mathfrak{I}}(\theta)}]$$

*O(log n / n)*


- Where

$$\mu = 144 \ln m \ln(2|H|)/\theta^2 + \ln(2|H|/\delta)$$

$$\hat{\mathfrak{I}}(\theta) = \Pr_S[yf(x) < \theta] \Pr_S[yf(x) \geq \frac{2\theta}{3}]$$

Related to margin variance

# Experimental results



➤ Margin variance really important

Figure from [Gao & Zhou, AIJ 2013]

Figure 1: Each curve represents a voting classifier. The  $X$ -axis and  $Y$ -axis denote example and margin, respectively, and uniform distribution is assumed on the example space. The voting classifiers  $h_1$ ,  $h_2$  and  $h_3$  have the same average margin but with different generalization error rates:  $1/2$ ,  $1/3$  and  $0$ .

[Shivaswamy & Jebara, NIPS 2011] tried to design new boosting algorithms by maximizing average margin and minimizing margin variance simultaneously, and the results are encouraging

# Summary

- ❖ AdaBoost is an effective ensemble strategy to combines multiple weak models into a strong one
  - Additive structure + focusing on mistakes
  - It's usually easy to find a bunch of weak models but hard to find one single strong model
  - The idea behind the algorithm is more important than the algorithm itself
- ❖ We can view the AdaBoost from different perspectives
  - It use a set up of forward stagewise additive modeling
  - It can be considered as a Newton method
  - The concept of margin explains why AdaBoost seems never overtrained

# Homework

For a  $K$ -class classification problem, we can recode the class label  $c$  with a  $K$ -dimensional vector  $\mathbf{y}$  with all entries equal to  $-\frac{1}{K-1}$  except a 1 in position  $k$  if  $c = k$ , i.e.,

$$y_k = \begin{cases} 1, & \text{if } c = k, \\ -\frac{1}{K-1}, & \text{if } c \neq k. \end{cases}$$

Let  $\mathbf{f} = (f_1(\mathbf{x}), \dots, f_K(\mathbf{x}))^T$  with  $\sum_{k=1}^K f_k(\mathbf{x}) = 0$ , and define

$$L(\mathbf{y}, \mathbf{f}(\mathbf{x})) = \exp\left(-\frac{1}{K} \mathbf{y}^T \mathbf{f}(\mathbf{x})\right).$$

- (a) Using Lagrange multipliers, derive the population minimizer  $\mathbf{f}^*$  of  $\mathbb{E}_{\mathbf{y}|\mathbf{x}}[L(\mathbf{y}, \mathbf{f}(\mathbf{x}))]$ , subject to  $\sum_{k=1}^K f_k(\mathbf{x}) = 0$ , and relate these to the class probabilities.
- (b) Derive a multiclass boosting algorithm using this loss function and verify that it covers the Adaboost algorithm as a special case ( $K = 2$ ).
- (c) Implement your derived algorithm, where you are allowed to call package of trees. Compare your implementation with the existing standard gradient boosting package on a multiclass classification problem. Make some discussion about what you observed. (Hint: there are two key differences: First, unlike adaboost which uses classification trees, gradient boosting uses regression trees. Second, gradient boosting has a shrinkage parameter ( $\lambda$ ), a smaller learning rate ( $\lambda = 0.01$ ) often dramatically improves the performance compared with  $\lambda = 1$ ).



# Backups



# Reference

- ❖ ISL, ESL
- ❖ [Boosting 25年 \(2014 周志华\) - Bilibili](#)
- ❖ The Annals of Statistics 2000, Vol. 28, No. 2, 337-407